

A Comparison of Algorithms for Building ST-Histograms

José A. Muñiz Navarro

MIT

December 10, 2008

1 Motivation

1 Previous Work

1 Proposal

Histograms help database systems optimization

- Tables have fields
- One histogram per field.
- Selectivity helpful for optimizing queries


Need to sort by age

- Insertion sort?
- Merge sort?

ID	Age
1	17
2	21
3	18
4	19
5	18
6	21
7	17


How to calculate selectivities from histograms

FIND students WITH AGE $\text{age} \leq 14$



How to calculate selectivities from histograms

FIND students WITH AGE age ≤ 14

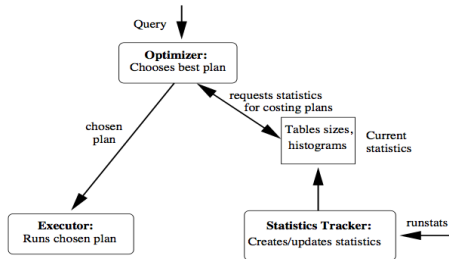


$$\sigma = 5 + \frac{10}{2} = 10$$

Building histograms: Traditional cost-based optimization

Four components:

- Optimizer
- Executor
- Histogram
- Statistics gatherer



No connection between executor and histogram.

Some problems with traditional approach


Problems:

- Tradeoffs
 - performance \leftrightarrow adaptability
 - performance \leftrightarrow precision
- Postgres suggests turning off statistics analyzer for large tables!

Idea: Incremental build via Self Tuning Histograms

Previous work by


- Aboulnaga and Chaudhuri - ST Histograms
- Babu, Bizarro - Adaptive Query Processing



Idea: Incremental build via Self Tuning Histograms

Previous work by

- Aboulnaga and Chaudhuri - ST Histograms
- Babu, Bizarro - Adaptive Query Processing



Interface


- Executor executes
`emit([a...b], val)`
- Histogram builder provides
`int estimate([a...b])`

Self Tuning Histograms: Interface and Motivation




- How do we build the histogram without looking at data?
 - Update frequencies
 - Update buckets

Updating frequencies

Result from `emit([0,19], 30)`



Merge and split



Proposed work


Determine feasibility of implementing ST Histograms in transactional databases, by benchmarking different ST algorithms against a common database system.

Algorithm modifications:

- Occasionally renormalize data to fit to total number of tuples.
Improves determination of error vs insertion
- Adjust blame proportional to range and frequency
- Split most frequently updated ranges

Proposed Setup

- Middle layer between database and end user
- Process:
 - 1 Relay query to database
 - 2 Obtain query plan from database
 - 3 Simulate query, using results for calling emit
 - 4 Calculate error rate in generated histograms
- Measure performance gains over common statistics gathering in transactional loads.



Expected Results

- Fast to adapt
- Minimal insertion overhead (Small amount of tuples per query)
 - Ideal for large databases with constant insertions and varying ranges.

Caveats:

- Major optimizer misses due to independence assumptions, and not failures in row estimates.
- Not all cases lead to performance gains.
- No solution for LIKE queries.

- Week 1: Build parser
- Week 2: Build executor
- Week 3: Build histogram builder
- Week 4-5: Implement different histogram building algorithms
- Week 6-7: Implement TPC-C benchmark
- Week 8-10: Run statistics and gather data